

TEXT-ONLY PITCH DECK

shhMOOD

The moodboard that writes the brief.

An AI canvas that turns image boards into cohesive prompts and text boards into reusable skill.md voice profiles.

Functional React MVP.

Image boards -> prompts. Text boards -> skill.md.

PROBLEM

Creative references pile up. Briefs still get rewritten.

Moodboards show taste, but they do not explain it well enough for models, teammates, or repeatable production.

WHERE THE WORK BREAKS

- Visual intent stays vague: palette, lighting, texture, framing, and atmosphere remain open to interpretation.
- Writing voice is hard to transfer: samples reveal rhythm and diction, but rarely become a reusable guide.
- Prompts get rebuilt from scratch whenever the model, project, or collaborator changes.
- The missing layer is synthesis. References should become instructions.

MARKET TIMING

Why now

Teams already use AI, but the handoff from taste to usable direction is still manual.

SIGNALS

- McKinsey 2025: 88% of surveyed organizations report regular AI use in at least one business function, but only about one-third have begun to scale AI programs.
- Adobe 2025: nearly two-thirds of marketers expect content demand to quintuple between 2024 and 2026.
- Canva/Harris Poll 2025: 84% say poor visual communication causes delays and confusion; the U.S. visual content economy is estimated at \$143B annually.
- Adobe Creators Toolkit 2025: 86% of creators use creative generative AI, and 85% would turn to an AI agent that learns their creative style.

GAP

- Generation is getting easier. Precise creative direction is still hard to capture, reuse, and share.

SOLUTION

shhMOOD makes references operational.

Collect on an infinite canvas, then turn the board into an output that travels.

TWO BOARD TYPES

- Image board: photos -> image analyses -> one cohesive image-to-text prompt.
- Text board: writing samples -> synthesized skill.md voice profile.

RESULT

- The output can be copied, downloaded, regenerated, or reused in another model or workflow.
- The product shifts moodboards from "look at this" to "make more like this."

HOW IT WORKS

The product loop

Boards, canvas, and output stay in one workspace so the creative loop stays tight.

FLOW

- Create: click New, name the board, and choose Image or Text.
- Collect: drop images, text files, selected text, or typed notes.
- Arrange: pan, zoom, and move items freely on an infinite canvas.
- Synthesize: output updates when content changes, not when layout changes.
- Reuse: export the prompt or skill.md as the handoff asset.

IMAGE BOARD

Image boards turn visual taste into one prompt.

Each photo is read, then the set is fused into a single creative direction.

WHAT HAPPENS

- Each image is downscaled and analyzed for subject, composition, color, texture, lighting, atmosphere, and style.
- The analyses are synthesized into one prompt, not a list of separate references.
- Add or remove an image and the prompt re-fuses automatically.
- Moving items never triggers generation.

EXAMPLE OUTPUT

- A hushed winter cabin at blue hour, snow-muffled light, warm lamp glow, pine, woodsmoke, indigo and amber.

TEXT BOARD

Text boards turn voice into skill.md.

Writing samples become a reusable Markdown style profile once the board has enough material.

HOW IT WORKS

- Drop text files, paste selected text, or add notes directly on the canvas.
- Below five notes, the app shows progress toward generation.
- At five or more notes, it creates a skill.md covering voice, tone, rhythm, diction, punctuation, themes, structure, do/don't rules, and a voice cheat-sheet.
- Edits update the profile, so the skill stays alive as the board changes.

USE

- A portable voice guide for writers, editors, agents, custom GPTs, and model workflows.

PRODUCT BEHAVIOR

The live loop stays fast and predictable.

The app separates content changes from layout changes, so synthesis only runs when it should.

RELIABILITY DETAILS

- Content edits are debounced by 650 ms before output generation.
- Async generations are token-guarded; stale results are discarded after rapid changes.
- Pan, zoom, and drag do not trigger model calls.
- Errors surface with a regenerate path instead of a silent hang.
- The current MVP stores work in session memory, so a refresh starts clean.

ARCHITECTURE

Built on a provider-aware completion layer.

The workflow is the product. The model provider can change.

TODAY

- Default path uses Claude/Anthropic inside the artifact preview.
- Image boards send images plus instructions. Text boards send samples plus skill.md instructions.
- One completion function routes the request based on provider and content type.

NEXT PRODUCTION STEP

- Move provider calls behind a small backend so API keys are not exposed in the browser.
- Keep prompts and Markdown skills model-agnostic so outputs remain portable.

USE CASES

For teams that need creative direction to travel.

shhMOOD is useful anywhere examples need to become repeatable instructions.

PRIMARY USERS

- Creative directors: turn visual references into prompts for concepts, campaigns, and art direction.
- Brand teams: turn visual and verbal systems into repeatable briefing assets.
- Writers and editors: convert samples into a practical style guide for drafting and review.
- AI builders: generate prompt and skill assets for agents, evaluations, and custom workflows.
- Agencies and studios: reduce the translation gap between inspiration, brief, and production.

DIFFERENTIATION

Not a static board. A synthesis layer.

The core product is not storage. It is conversion from references to reusable direction.

EDGE

- Static moodboards store references. shhMOOD converts references into instructions.
- Prompt tools start from text. shhMOOD starts from actual image boards and writing samples.
- Manual style guides are slow. shhMOOD creates a working skill.md from examples.
- The canvas supports spatial thinking while the output panel stays live.
- The deliverable is the prompt, the style guide, or the brief - not just the board.

BUSINESS MODEL AND GTM

A narrow wedge with room to expand.

Start with creators and teams already building moodboards, prompt packs, style guides, and AI workflows.

INITIAL PATH

- Private beta with creators, brand teams, agencies, and AI workflow builders.
- Freemium single-user canvas; paid Pro for saved workspaces, exports, and provider configuration.
- Team tier for shared libraries, version history, brand/style systems, role access, and admin controls.
- Expansion into prompt testing, creative-brief export packs, design-tool integrations, and enterprise privacy controls.
- Core metric: reusable outputs created and reused from boards.

ROADMAP AND ASK

Built now. Clear next steps.

The MVP proves the loop: collect examples, synthesize them, and export reusable creative direction.

MVP TODAY

- Functional React app with image boards, text boards, infinite canvas, real model synthesis, copy, download, regenerate, and error states.

NEXT BUILD MILESTONES

- Persistent workspaces, shareable boards, version history, team libraries, prompt testing, and export packs.

ASK

- Recruit beta users who already use moodboards, prompts, style guides, or AI creative workflows.
- Target proof: users turn a messy reference set into a prompt or skill.md in minutes, then reuse it in another tool.

SOURCES

References

Text-only deck sources used for product framing and market context.

PRODUCT SOURCES

- User brief and uploaded shhMOOD React app.
- shhMOOD 30s spot: "the moodboard that writes the brief" and image boards -> prompts / text boards -> skill.md positioning.

MARKET SOURCES

- McKinsey, The State of AI: Global Survey 2025.
- Canva, Visual Communications Report 2025, with The Harris Poll and Neuro-Insight.
- Adobe, GenStudio Content Supply Chain press release, March 2025.
- Adobe, Creators Toolkit Report, October 2025.